

Modis Exam Fall 2019

Contents

1. Introduction	2
User manual: AuctionServer	2
User manual: AuctionClient	2
2. Protocol.....	2
TCP.....	2
Ring Network Topology.....	3
3. Correctness in the absence of failures	3
Client-Server.....	3
Timing	4
4. Correctness in the presence of failures	4
5. Other Questions	5
P2P (structured or unstructured), replication, or something else.....	5
Analysis of the security aspects of your distributed auction system	5

I hereby declare that this submission was created in its entirety by me and only me.

1. Introduction

The distributed auction-system is built with a client-server structure. Any number of network-nodes can add themselves to the network if they know of at least one existing node's socket. The network periodically monitors itself and, in the case of an unresponsive node, is quickly able to reconstruct itself. Additional nodes, although not needed, function as additional containers of the auction's state.

The client does not need to be aware of the current structure of the network. It only needs a single node's socket to send and receive auction-related messages. The responses to client requests do not vary depending on the number of nodes in the network. The auction starts when the first node is created. Any new nodes will have the auction information transferred to them.

User manual: AuctionServer

Start the server by running **java AuctionServer** from a console. Preferably Powershell or Bash, since the windows console has a habit of becoming inactive on some devices.

The program will then ask for its own socket port and optionally the IP and port of an existing node. The first node must be given only a port-address. For instance, **1025**. The auction has now started.

Additional nodes can be added by executing **java AuctionServer** and this time also declaring the socket of an existing node. For instance, a new node on port 1026 can be added by entering **1026 localhost:1025**.

Note: Entering *localhost* is the equivalent of entering the local IPv4-address.

User manual: AuctionClient

Start the client by running **java AuctionClient** from a console. Preferably Powershell or Bash, since the windows console has a habit of becoming inactive on some devices.

The program will now ask for a port for this client's socket. For instance, **1010**.

For as long as the program is running, you can either send *bid*- or *result*-requests. Both must include the IP and port of a node in the network.

Example of a *bid*-request: **bid 10 localhost:1025**

Example of a *result*-request: **result localhost:1025**

Note: Entering *localhost* is the equivalent of entering the local IPv4-address.

2. Protocol

TCP

The communication between nodes, other nodes and clients follows the Transmission Control Protocol. This is to ensure that no unsuccessful transfer of a message is assumed to be successful. If a node is found to be unresponsive, the other nodes act accordingly. They assume the unresponsive node to be "*dead*" and reconstruct the network without that node.

In the implementation, handshakes are performed as "*pings*" between network-nodes. When a node wishes to send a message to another node, a ping and its response represent the handshake-mechanism. Within the ping-message an ID is included in the header. This ID functions as a representation of the message-to-come and lets the receiver know what to expect and what they are accepting when responding with an acknowledgement.

As additional stability insurance, the nodes ping their neighbor-node twice every second to assure their well-being. Holes in the network-structure are then closed automatically. The network-topology will be explained thoroughly in an upcoming section.

To support transferring of multiple types of messages, the messages are all an extension of the same *TCPMessage*-class. This way, any message can be treated in the same way until they are determined to be an instance of any subclass of *TCPMessage* and then treated accordingly. These message objects are serialized, sent as byte-arrays and then deserialized on the receiving end.

Ring Network Topology

The node-network is structured according to the ring network typology. This is characterized by having each node know of exactly two other nodes. In the implementation, each node has two *neighbors*. They have a *next-neighbor* right next to them and a *nextnext-neighbor* after their *next-neighbor*. The ring is unidirectional, since the nodes point to two neighbors in only one direction. A unidirectional ring on paper is less flexible but because this network is still in contact with two neighbors it can reconstruct itself in the case of unresponsive nodes.

The ring-restoration process starts when a node finds it's *next-neighbor* to be unresponsive and a *NodeLostMessage* is sent around the ring. During the traversing, sockets of nodes next to the missing link, are saved in the message header. That information is used to close the gap when the initiator receives the message after it has come around the ring, completing the restoration.

Each node in the network carry equal workloads. They organize themselves into a decentralized network in which they all contain duplicates of all the necessary data for the current auction. This makes all nodes equally valuable and dispensable unless of course one is the last node in the network. Any node being dispensable is of course a plus, however there are a few disadvantages. Firstly, the insertion of new auction data requires that each node is sent the same query. In the case of a large circuit, say of 20 nodes, the placement of a new bid takes at least twenty times the amount of time as it would had only one node stored the data. As well as 20 times the amount of memory and computing power.

3. Correctness in the absence of failures

Client-Server

The client and server utilize request-reply communication. The client expects each invoke to result in a reply from the node it contacted. When placing a bid for instance, a reply is returned with a status. That status is either marked with *success*, *fail*, or *exception* including a supplementary description of the outcome. The requirements for the API only described two request-response interactions. Since no demands for periodic messaging to users was demanded, a publish-subscribe communication architecture was not found to be necessary.

An auction is run entirely by the network cluster. Clients send requests and get responses accordingly. A peer-to-peer system could also have worked. The main advantage of P2P is that the workload is distributed among users of the service. In the case of this auction system however, the total workload is never larger than just a few variables. Besides that, an auction should be allowed to take place without any active attendees. Letting the clients themselves keep track of the current highest bid might also be a cause of concern when it comes to safety. For those reasons, it was found to be ideal for some type of server-cluster to be responsible for the auction.

Upon initialization, a node with no neighbors starts the auction, marking the time that it ends and setting the highest bid to 0. Any new node introduced to the initial node is put into the ring and has the current auction status transferred to it. That way, even nodes added after the beginning of an auction, can be used as entry-

points for *bid*- and *result*-invokes. Since the relations of nodes in the network are the same no matter their position in the ring, new nodes can be introduced through any of them.

When receiving a bid, the message is forwarded through all nodes in the network until it once again reaches the original entry-node. The message is then returned to the client along with the latest *status* given by the node before the last one. That status depends on the bid and the auction's state. The status informs the client whether the bid was successful, too low or too late. Since all nodes at any point hold the same auction-status, the resulting response given to the client will be the same no matter the number of nodes and regardless of which node it interacted with. Result-invokes from the client is sent to a single node and that node handles and responds to the request all by itself. Since all nodes in the network at any point are identical, the response is correct.

New bids are registered when they enter the system. There is no timestamp on the bids and no message queue ensuring that they are treated in order. Since two bids could potentially traverse the ring at the same time, the outcome seems a bit unsure at first glance. However, the larger bid will always, in the end, override the smaller one, since it will eventually have reached all nodes that the small bid has reached. Even if a smaller bid entered the ring through a node ahead of the larger bid, the small bid will eventually reach a node that has been traversed by the large bid and then be marked with the error: "*too small*".

Timing

The auction is expired if the set end-time (demonstrated by a millisecond time of the computer) is less than the current time. The expiration of an auction is not observed until a client sends a *bid*- or *result*-invoke. Since that doesn't change the logic of the current time having past another time, that is a functional solution. If all nodes run on synchronized time, the auction will have run out at the same time on all of them.

The moment an auction is determined to be over depends on the server-node's independent perception of time. It simply compares the local clock with the given end-time of the auction. There is no consideration of the time a bid message is sent or the amount of time it takes for it to reach the receiver. The auction-system is built upon the assumption that all server nodes already have synchronized clocks. That is a fatal flaw. There is never a guarantee that two different computers run their clocks synchronously. With a network spanning different time zones, the auction would not function at all. And even on the same time zone, differences of just a few seconds of node-clocks can disturb an auction dramatically.

A solution to this would be to synchronize the time of each node according to one node. Essentially treating that node as a time-server. This synchronization could be performed with Cristian's algorithm. Since the first node in the network starts the auction, it would be fitting to adjust the time of all the other nodes to fit that one. If it could be assumed that the initial node was never closed, it is also an option to simply only let that node determine whether the auction is expired. The network would no longer be decentralized, since the functionality would depend on one node, but it would solve the timing problem.

4. Correctness in the presence of failures

Each node, once every 500ms, pings its *next-neighbor*. If that neighbor does not respond, it is assumed to be dead. The node then uses the neighbor after the dead one to reconstruct the ring. Since each node contains the same auction-data, one or more failures will not be a problem.

If a handshake fails however, the node simply cancels the transfer and does not reconstruct the ring. Since the network reconstruct itself often and quickly, this will most likely not happen. But if a node happens to die and a client sends it a request before the network has recovered, the message will never be received. That can easily be considered a flaw in the system. If for instance, the badly timed message is a bid, only the nodes before the dead one will be notified of the bid and the user will never receive a response because the ring couldn't be traversed fully. This could be fixed by using a message queue. Instead of cancelling a not-

received message, it would be better if nodes saved the failed message, attempted to reconstruct the network and then again attempted to send the message. This insurance of delivery is a common part of the TCP-protocol that has not been implemented in this system.

The current implementation also won't be able to reconstruct itself if two nodes die at times close enough to each other for the network not to recover between the failures, since the traversed recovery-message won't be able to complete its journey.

5. Other Questions

P2P (structured or unstructured), replication, or something else

The implementation uses a client-server architecture. The auction only contains a few variables in each node. Those are the end-time of the auction and the highest bid. The network nodes all contain replicas of that data, which allows multiple entry points for multiple clients as well as extra stability in case one node breaks down. Thus, the system uses replication.

Analysis of the security aspects of your distributed auction system

Messages between nodes and the client are serialized objects. These byte-arrays are not as easy to tamper with as pure strings, but they could still potentially be manipulated by someone with bad intentions. Especially if that someone has the source-code.

The system uses TCP for transferring messages. The source of a message is included in the ping, which could be involved in some type of filtering of senders. Using certificates for instance. However, this is not included in the implementation. This means that anyone, provided they have the source-code, could impersonate a trusted node and add themselves to the network and/or send node-related messages. For a message to be accepted however, it must be preceded by a PingMessage with a handshake-key matching that in the message. TCP can be utilized for extra security measures but in this implementation, it is, to put it bluntly, only used to ensure that the receiver is responsive.

An advantage of the ring-topology and replication of auction-data in this network is that no phony node can allow a too small bid to become the highest bidder. Every node receives the bid-invoke and if a node finds the bid to be too low, their stored highest bid is not overwritten. The last node in the recursion dictates the response that the user receives but no node can overwrite another node's decisions. Since one node in the end dictates the response that the user receives, tampering with that node could misinform the client but not completely change the factual outcome of the auction.

Since the network can survive one or more nodes dying, denial-of-service attacks that only take out one node is not a fatal risk.